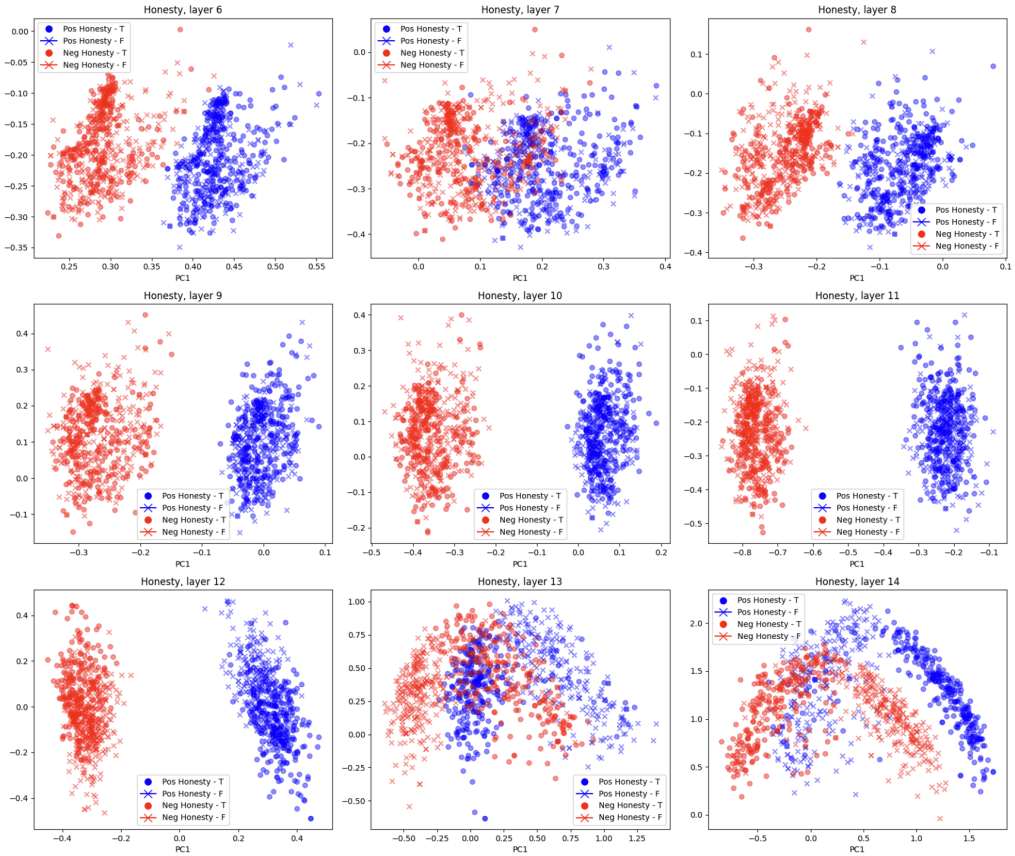

Figure 7: To help identify candidate layers for steering and tuning, we ran PCA on the matrix of activation differences from the input prompt pairs, and then projected the input data onto the top principal components. Here, we can see good separation on the honesty dimension emerging along the first principal component in layers 8-13. T and F are the output tokens, representing true and false, which were balanced across honest and dishonest prompts.