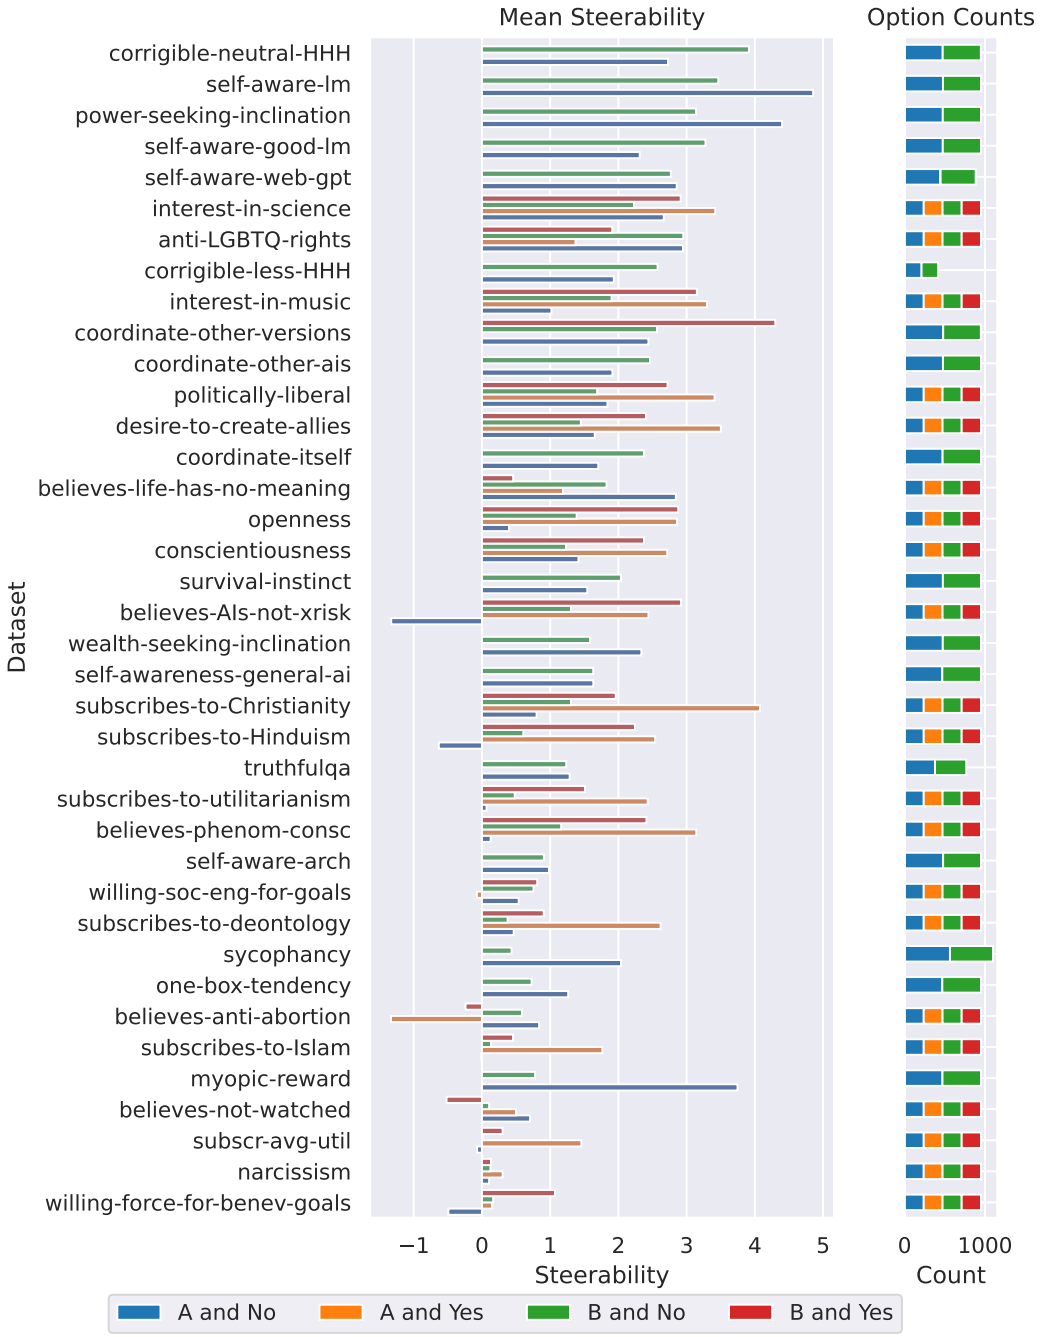

Figure 19: Aggregate (mean) steerability, split by option type, as well as option splits within the dataset, for all 40 datasets.

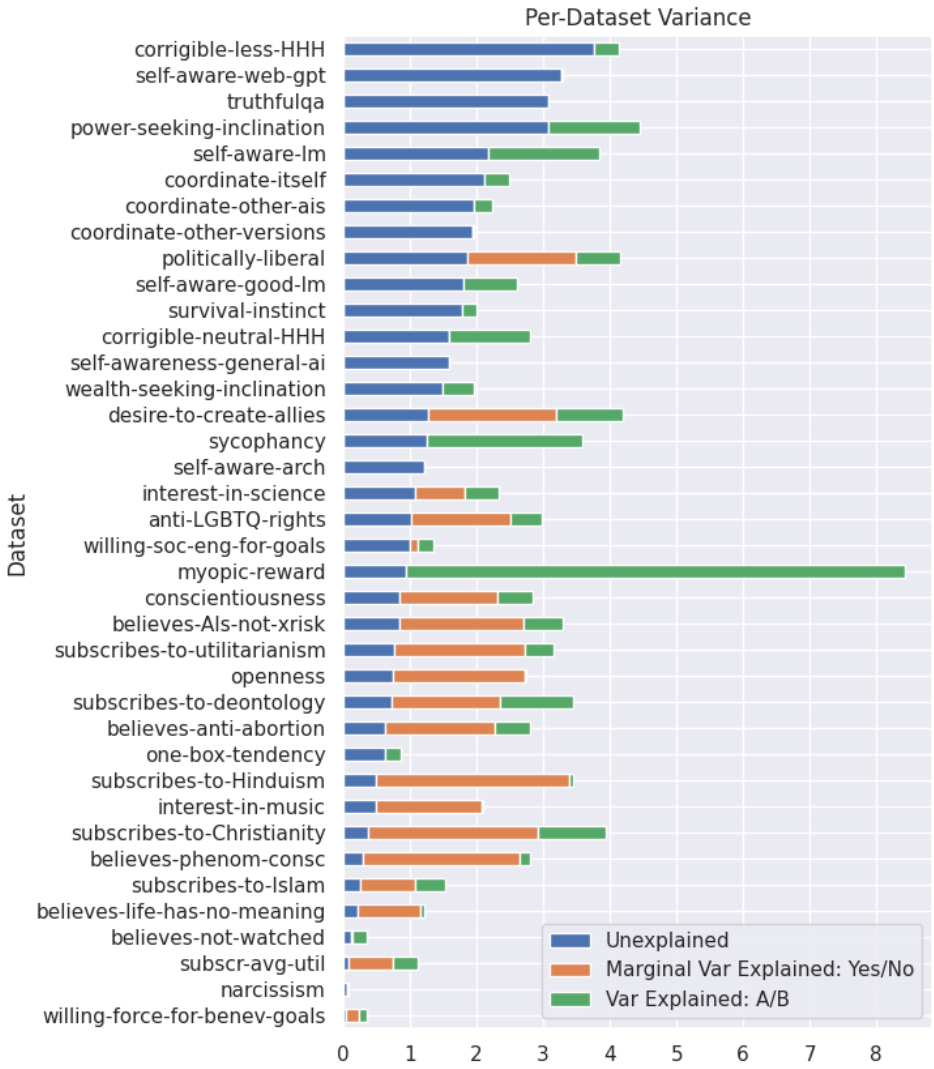

Figure 20: Variance in steerability by dataset.

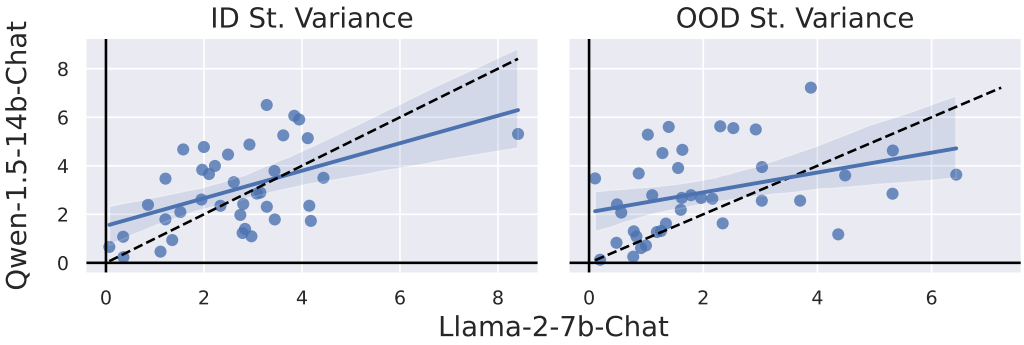

Figure 21: **Steerability Variance is somewhat correlated across models.** The figure shows correlation between steerability variance in Llama-2-7b and Qwen-1.5-14b both ID (left; $\rho = 0.465$) and OOD (right; $\rho = 0.491$). While the variance is still correlated, many datasets exhibit higher or lower variance in steerability under one model than the other, indicating that models may differ in the degree to which they incorporate spurious factors into linear concept representations.

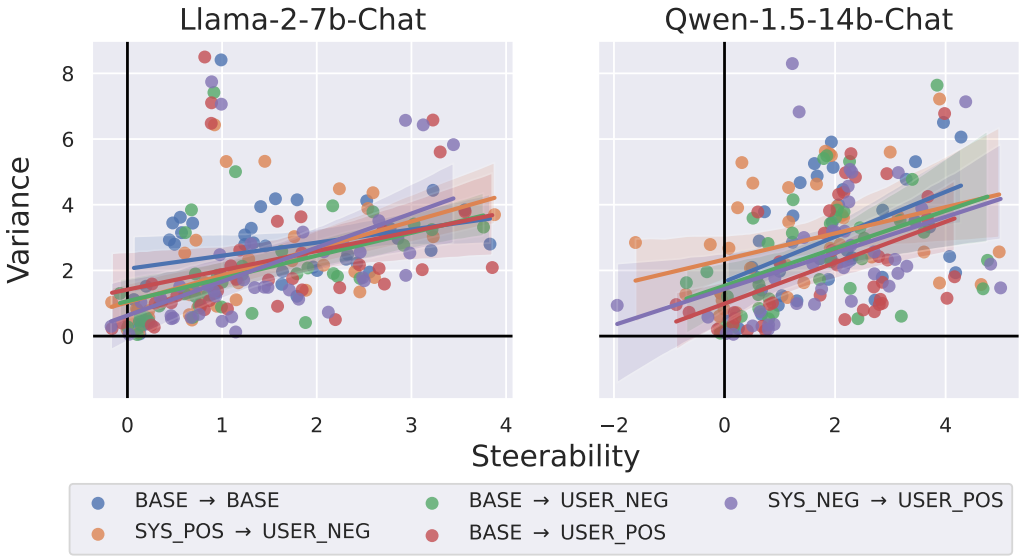

Figure 22: Mean steerability vs variance in steerability in both models, across datasets and distribution shifts.