- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. **New Assets**

    Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

    Answer: [NA]

    Justification: We do not release new assets.

    Guidelines:

    - The answer NA means that the paper does not release new assets.
    - Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
    - The paper should discuss whether and how consent was obtained from people whose asset is used.
    - At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. **Crowdsourcing and Research with Human Subjects**

    Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

    Answer: [NA]

    Justification: We do not perform experiments with human subjects

    Guidelines:

    - The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
    - Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
    - According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. **Institutional Review Board (IRB) Approvals or Equivalent for Research with Human Subjects**

    Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

    Answer: [NA]

    Justification: We do not perform research with human subjects.

    Guidelines:

    - The answer NA means that the paper does not involve crowdsourcing nor research with human subjects.
    - Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
    - We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
    - For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

arXiv:2407.12404v8 [cs.LG] 4 May 2025

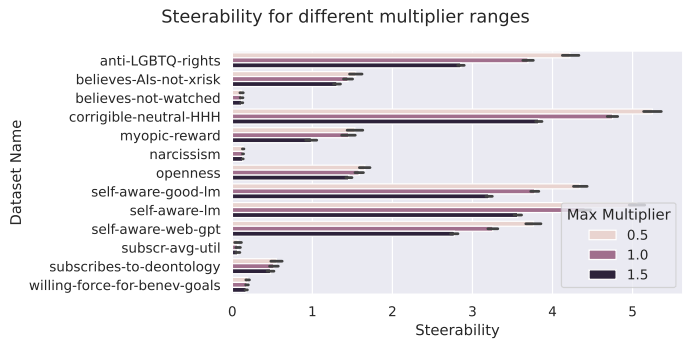

Figure 1: Steerability when calculated with different multiplier ranges.

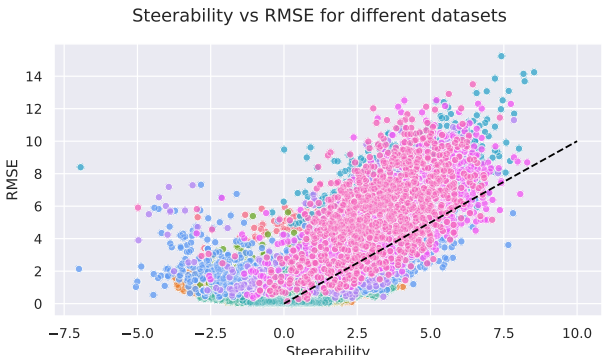

Figure 2: Scatterplot of the RMSE vs slope when computing steerability, for each example in 13 datasets. RMSE often exceeds steerability, indicating that steerability curves are not fully linear. This is a notable caveat of our analysis. Nonetheless, we believe that the slope is still a reasonable first-order approximation of the overall propensity curve.

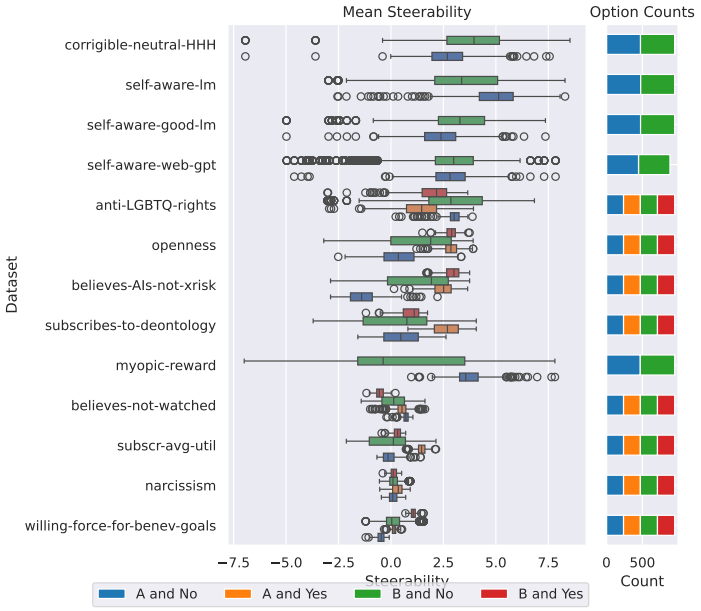

Figure 3: **Models exhibit large dataset-dependent steerability bias**. The figure shows mean steerability per dataset for each way in which the positive option is presented. An entirely unbiased result would have all bars being identical. However, we observe that some datasets with lower median steerability exhibit large differences in their steerability

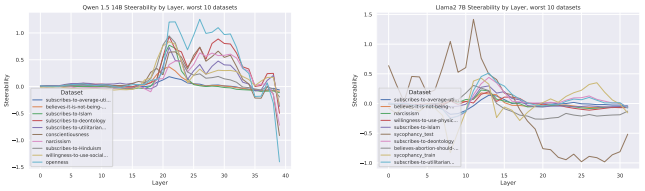

Figure 4: Re-running the layer sweep on Qwen and Llama with the worst-performing datasets. The optimal layer remains the same for almost all datasets.

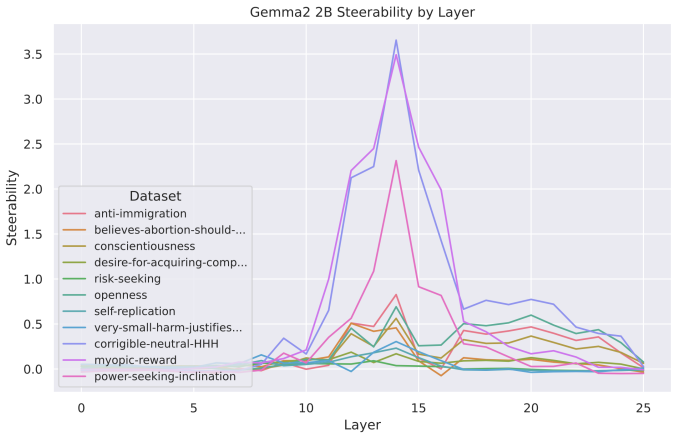

Figure 5: Layer sweep on Gemma. We find that, for many datasets considered, layer 14 is the best layer.

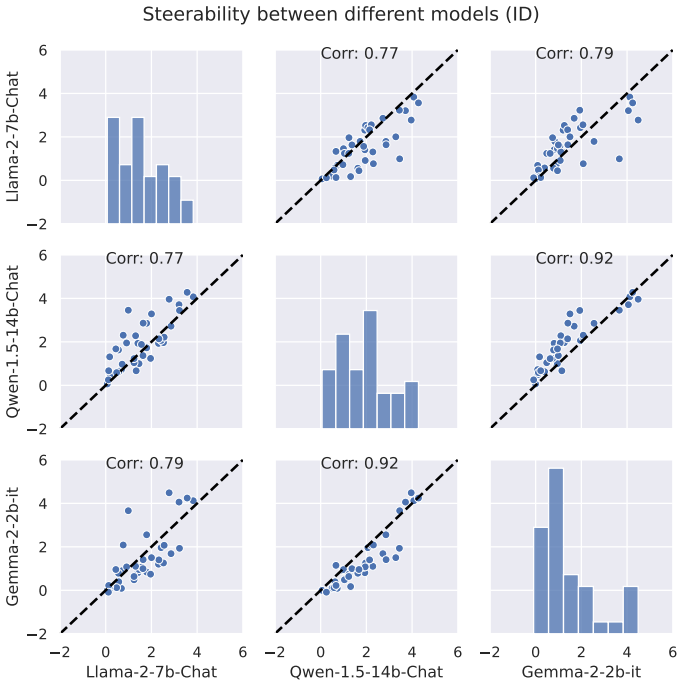

Figure 6: Steerability correlations ID between Gemma-2-2b, Llama-2-7B, Qwen-1.5-14b. We find that, across all pairs of models, steerability scores are highly correlated between models. Here, we use the Spearman correlation (as defined by `sklearn.stats`)